

FaithForge Technologies & Consulting LLC

AI API Cost Proposal - GPT-4o-mini

Monthly operating cost breakdown for the FaithForge Contract Opportunity Screener

Date: June 16, 2026

Classification: Internal - Budget Approval

EXECUTIVE SUMMARY

FaithForge operates an internal AI-powered contract screening platform that automatically scans incoming emails for RFP/RFQ/grant opportunities, reviews solicitation documents, and generates submission-ready proposal packets. The system currently runs on the Groq free tier, which imposes hard daily token limits that are consistently exceeded during normal operations.

This document requests authorization to subscribe to the OpenAI GPT-4o-mini API (pay-as-you-go, no minimum commitment). Total projected cost is \$2.62 per month based on current usage patterns.

\$2.62 / month

Total projected API cost | \$31.41 annually | No minimum commitment | Cancel anytime

CURRENT PROBLEM - GROQ FREE TIER LIMITATIONS

FREE TIER CONSTRAINTS

- 100,000 token daily hard cap
- 12,000 tokens/minute rate limit
- Scanning 30 emails alone consumes 85,000 tokens - leaving almost nothing for packet builds
- Packet builder takes 3-4 minutes due to forced retry waits on rate limit errors
- Daily quota exhausted by mid-day, blocking all AI operations for the rest of the day

WITH GPT-4o-mini PAID

- No daily token cap
- 200,000 tokens/minute rate limit
- Packet builds complete in ~30 seconds (vs 3-4 minutes)
- All operations run without interruption throughout the day
- Pay only for actual usage - no wasted quota

AI OPERATIONS BREAKDOWN

The platform uses AI for three distinct operations. Token consumption was measured directly from application code.

Operation	What It Does	Calls/Run	Tokens/Run	Daily Volume
Proposal Packet Builder	Generates full submission-ready proposal (Executive Si 6		28,511	5 builds/day
Email Screener	Classifies emails as Relevant/Possibly/Not Relevant; e> 1		3,328	~30 emails/day
Document Reviewer	Reads uploaded solicitation docs; extracts eligibility, for 1		4,570	~3 reviews/day

MONTHLY COST DETAIL

Pricing: GPT-4o-mini - \$0.15 per 1M input tokens / \$0.60 per 1M output tokens (verified from OpenAI pricing page, June 2026).

Operation	Monthly Volume	Input Tokens	Output Tokens	Input Cost	Output Cost	Subtotal
Packet Builder	150 builds	1.70M	2.58M	\$0.25	\$1.55	\$1.80
Email Screener	900 emails	2.46M	0.54M	\$0.37	\$0.32	\$0.69
Document Reviewer	90 reviews	0.28M	0.14M	\$0.04	\$0.08	\$0.12
TOTAL MONTHLY	-	4.44M	3.26M	\$0.66	\$1.95	\$2.62

USAGE SCENARIO ANALYSIS

Projected cost under different usage intensities. Even at maximum conceivable usage the cost remains under \$10/month.

Scenario	Emails/Day	Doc Reviews/Day	Packet Builds/Day	Monthly Cost
Current / Projected (Recommended)	30	3	5	\$2.62
Heavy Usage	100	5	5	\$5.22
Maximum / Worst Case	200	10	5	\$8.11

ALTERNATIVES CONSIDERED

Option	Model	Monthly Cost	Quality	Notes
RECOMMENDED	GPT-4o-mini	\$2.62	Very Good	Best price-to-quality ratio for proposal writing
Option B	Groq Dev Tier (llama-3.3-70b)	\$3.25	Good	No code changes; still has per-minute rate limits
Option C	GPT-5.4-nano	\$6.30	Good	Newer but nano-tier; 2.4x more due to output pricing
Option D	Claude Haiku 4.5	\$14.50	Best	Highest quality; output cost makes it expensive
Option E	Groq Free Tier (current)	\$0.00	Good	Daily quota exhausted mid-day; operations blocked

RETURN ON INVESTMENT CONTEXT

\$2.62

Monthly API cost

\$31.41

Annual API cost

1 contract

Covers 10+ years of API

A single awarded contract (typically \$50,000-\$500,000 for government PMO engagements) covers the entire annual API cost by a factor of 1,500x-15,000x. The \$31.41 annual cost is negligible relative to the business development value delivered.

SETUP & BILLING

No subscription or minimum commitment. OpenAI's API is billed strictly on usage. Setup requires:

- 1. Create an OpenAI account at platform.openai.com
- 2. Add a payment method and purchase a minimum \$5 credit
- 3. Generate an API key and provide it to the development team
- 4. Configuration change takes under 10 minutes - no downtime
- 5. Set a \$10/month spend cap in the OpenAI dashboard as a safeguard

A monthly spend cap can be configured in the OpenAI dashboard to ensure costs never exceed a defined ceiling regardless of usage volume.

RECOMMENDATION

Authorize purchase of an OpenAI API key (GPT-4o-mini, pay-as-you-go)

Initial credit: \$10 | Monthly spend cap: \$10 | Sufficient for ~4 months of normal operations

This eliminates all current free-tier limitations, reduces packet build time from 3-4 minutes to approximately 30 seconds, and ensures uninterrupted AI-assisted contract screening throughout the business day.